




University of Udine

Data Management for Big Data

Introduction to Big Data and NoSQL

Andrea Brunello

andrea.brunello@uniud.it

Big Data



Big data are beyond the usual limits of traditional databases, and are characterized by one or more of the properties:

- huge **Volume**
- high **Variety**
- acquired at high **Velocity**

Gartner analyst Doug Laney introduced the 3Vs concept in a 2001 MetaGroup research publication: "*3D data management: Controlling data volume, variety and velocity*"



Volume basically means the size of stored data, which may derive from human actions or can be machine-generated.

Sometimes **the volume of data is so massive** that they cannot be stored in their entirety, but have to be compressed/transformed online, as soon as they arrive (e.g., scientific sensor data).

Sometimes data can be stored using traditional RDBMS, while other times this choice may end up being too expensive in terms of cost or time

- ⇒ NoSQL solutions for the storage
- ⇒ frameworks like Hadoop for the processing



Volume (Orders Of Magnitude)

IBM 350 disk storage
(1956, 3.75 MB)

Walmart's DW
(1992, 1 TB)

1 year of CERN's
LHC data (15 PB)

Quantities of bytes						
Common prefix				Binary prefix		
Name	Symbol	Decimal	Binary	Name	Symbol	Binary
		SI	JEDEC			IEC
kilobyte	KB/kB	10^3	2^{10}	kibibyte	KiB	2^{10}
megabyte	MB	10^6	2^{20}	mebibyte	MiB	2^{20}
gigabyte	GB	10^9	2^{30}	gibibyte	GiB	2^{30}
terabyte	TB	10^{12}	2^{40}	tebibyte	TiB	2^{40}
petabyte	PB	10^{15}	2^{50}	pebibyte	PiB	2^{50}
exabyte	EB	10^{18}	2^{60}	exbibyte	EiB	2^{60}
zettabyte	ZB	10^{21}	2^{70}	zebibyte	ZiB	2^{70}
yottabyte	YB	10^{24}	2^{80}	yobibyte	YiB	2^{80}

In 2025, it is estimated that 175 ZB of data have been generated on the Internet:
https://en.wikipedia.org/wiki/Zettabyte_Era



Differences between formats and the absence of a common structure are a typical characteristic of big data.

Structured, semi-structured, and unstructured data.

Data may come from different sources. Considering the web it may come from humans, like *user-generated content*, or it can be machine-generated, such as *logs*, *packet traces*, etc.

Heterogeneity of formats, structures and sources make it difficult to process and store such data using traditional tools.




Data acquired via sensors, or scientific instruments, may come at a **high speed**.

Some data have to be stored or analyzed as soon as they arrive, since they are transient (e.g., logs, data streams).

For companies that rely upon fast-generated data it is also important to exploit/analyze such data as fast as possible.

“Just in its 1st phase, the SKA telescope will produce some 160 TB of raw data per second that the supercomputers will need to handle.” <https://www.skatelescope.org/frequently-asked-questions/>








Data *velocity* is rarely a static metric.

Internal and external changes to a system and to the context in which it operates can have a considerable impact on the rate at which data have to be managed.

Variable velocity, coupled with large volume, requires data stores to handle **sustained levels** of high read and write loads, and also **peaks**.





- **Volume:** scale of the data
- **Variety:** different forms of data
- **Velocity:** very high speed of the incoming data
- **Variability:** changes in the characteristics of the data
- **Value:** revenues, hypotheses that may arise from the data
- **Veracity:** trustworthiness, origin and reputation




- **7 Vs:** <https://impact.com/marketing-intelligence/7-vs-big-data/>
- **10 Vs:** <https://tdwi.org/articles/2017/02/08/10-vs-of-big-data.aspxs>
- **17 Vs:** <https://www.irjet.net/archives/V4/i9/IRJET-V4I957.pdf>
- **42 Vs:** <https://www.kdnuggets.com/2017/04/42-vs-big-data-data-science.html>



Big Data

Classification by Volume and Complexity





Data from the Web plays a big role in big data realm, and are characterized by volume, variety and velocity.


Examples:

- HTML pages
- Tweets
- Social network content (Facebook, LinkedIn, etc.)
- Forum comments and blog posts
- Documents in several formats: XML, PDF, Word, Excel, etc.



The Internet Of Things (IoT) is about (daily life) objects that are equipped with sensors and connectivity, acting as sources of data

- the concept can involve both industry and consumer products (like connected automobiles, machinery, etc.)
- data can be used for tasks such as surveillance, predictive maintenance, or performance enhancement in general
- if objects are attached to people even human behaviour and well-being can be analyzed





Big data come with challenges and opportunities:

- *business*: big data give companies the opportunity to develop new business models so to get advantages with respect to competitors
- *technology*: size and complexity of big data require adequate solutions
- *financial aspects*: several use cases show that exploiting big data may lead to economic benefits, as well as catastrophic failures. To this end, it is also important evaluate the costs involved with their management
- *privacy*: who owns the data and where is it geographically stored?

NoSQL



What is NoSQL?

The term NoSQL (Not only SQL) refers to data stores that are **not relational databases**, rather than explicitly describing what they are.

A possible (rather general) definition: *“Next Generation DBs mostly addressing some of the points: being non-relational, distributed, open source and horizontally scalable”*.

NoSQL storage technologies have very **heterogeneous** operational, functional, and architectural characteristics (*one size does not fit all*).

NoSQL proposals have been developed starting from 2009 trying to address new challenges that emerged in that decade.



Most enterprise level applications used to run on top of relational databases (MySQL, PostgreSQL, etc).

Over the years data got bigger in volume, started to change more rapidly, and to be in general more structurally varied than those commonly handled with traditional RDBMS.

As the *volume* of data increases dramatically, so does query execution time, as a consequence of the size of tables involved and of an increased number of JOIN operations (*join pain*).



High volume and high heterogeneity data:

- RDBMS: capacity and constraints at their limits
- NoSQL are specifically designed for this scenario

Elastic Scaling:


- RDBMS scale up: bigger load $\Rightarrow$ bigger server
- NoSQL scale out: distribute data across multiple nodes, adding/removing them seamlessly

Database Administrator Specialists:

- RDBMS require highly trained experts to monitor DB
- NoSQL require less management: automatic repair and simpler data models

Benefits of NoSQL

Vertical vs horizontal scaling



Scaling horizontally is not easy at all... we enter the realm of *distributed systems*!



The CAP Theorem

Also known as Brewer's theorem

When a network partition occurs, a distributed data store can guarantee at most **two** of the following properties:

- **Consistency (C).** Every read receives the result of the most recent successful write, or an error
- **Availability (A).** Every request receives a non-error response, even if the returned data does not reflect the latest written value
- **Partition tolerance (P).** The system continues operating despite network failures or communication breakdowns between nodes



In highly distributed systems, network partitions and node failures are frequent and unavoidable.

Therefore, **partition tolerance (P)** is typically non-negotiable.

When a partition occurs, systems must prioritize one of the following:

- **AP (Availability + Partition tolerance).**

The system remains available during partitions, but replicas may return outdated information or accept conflicting updates. Consistency is restored once communication resumes.

- **CP (Consistency + Partition tolerance).**

The system preserves consistency by rejecting or delaying requests when necessary, potentially sacrificing availability.



Several default constraints of RDBMS are not supported at the database level (e.g., foreign keys).

If not properly managed, the absence of a fixed structure may become an issue.

The design process is not as straightforward and consolidated as in the relational model.

Lack of SQL (though there are some SQL-inspired languages).


Often, lack of one or more ACID properties.



In order to better support the characteristics of the novel kinds of data, NoSQL systems rely on the BASE properties instead:

- **Basically available.** The store appears to be accessible most of the time (for instance, tolerance to node failures)
- **Soft state.** Stores do not need to be write-consistent, nor do different replicas have to be mutually consistent all the time
- **Eventual consistency.** The system eventually exhibits consistency at some later point


Of course, this kind of relaxed consistency would be a problem, for instance, in a banking database. In other cases, it is perfectly acceptable, e.g., social networks.



Some other technologies not included into the quadrant have also emerged, to cope with continuously evolving data management necessities and novel AI applications.

Most notably:

- Vector databases
- NewSQL





Key/Value Stores

Key/value stores are based on the concept of **associative array** (i.e., dictionary, or map), a data structure that contains couples of <key, values>; the key is used to access the values.

E.g., in a database storing personal data, one array could have as keys the SSN of customers associated to their phone numbers as values, while in another array the same keys could be associated with their profile pictures.

Operations allowed over an associative array are:

- *Add*: add a new pair to the array
- *Remove*: remove a pair from the array, based on the key
- *Modify*: change the value associated to a key
- *Find*: search for a value by the key



Array “emp_name”:

[<XA13, John Doe>, <XZ42, Mary Jane>, <XG56, Jhonny Donny>, ...]

Array “emp_department”:

[<XA13, Physics>, <XZ42, Physics>, <XG56, Computer Science>, ...]

Array “emp_salary”:

[<XA13, 45000>, <XZ42, 50000>, <XG56, 52000>, ...]

Array “dept_budget”:

[<Physics, 950000>, <Computer Science, 780000>, ...]

By convention, even if values are opaque and schemaless, all those included within a same array should be of a same kind.



Key/value stores offer just the add/remove/modify/find operations on data, which nevertheless are executed at a fixed (fast) computational cost, thanks to the associative array.

A value can be anything and of any size, but typically is totally opaque for the user (schemaless approach).

Some typical relational operations, such as complex filters and JOINS are not possible (unless a custom script is written).

Also, important RDBMS functionalities like foreign keys are not supported.

Key/value stores are very simple pertaining to their data model, but they can be extremely sophisticated regarding horizontal scalability.



Use cases:

- Storage of session data on the Web
- Storage of multimedia objects
- Storage of machinery or sensors data (large time series)

Exemplary DBMSs:

- **DynamoDB** developed by Amazon and available on the AWS (Amazon Web Services) platform
- **Project Voldemort** is a (LinkedIn-born) distributed key/value store based on Berkeley DB, designed for performance and fault protection
- **Redis** (in-memory database): <https://redis.io/>
- **Aerospike**: <https://aerospike.com/>



Document-oriented databases store, retrieve and manage document oriented information, also known as semi-structured data.

Documents do not have a fixed structure, nonetheless they contain **tags** or other markers to separate semantic elements and enforce hierarchies of records and fields within the data. Therefore, they can be considered as a kind of **self-describing structure**.

Documents can be encoded into formats like XML and JSON.


Document-oriented databases represent a step up with respect to key-value stores, defining a structure over keys and values, thus allowing the users to operate on the internal document structure.



Document-oriented Databases

Document example

```
<Phone campaign>
  <Survey ID = 1>
    <Name> John Easter </Name>
    <Question ID = 1> yes </Question>
    <Question ID = 2> yes </Question>
    <Question ID = 2.1> 5 </Question>
    <Question ID = 2.2> 7 </Question>
    <Question ID = 3> no </Question>
  </Survey>
  <Survey ID = 2>
    <Name> Mary Christmas </Name>
    <Question ID = 1> yes </Question>
    <Question ID = 2> no </Question>
    <Question ID = 3> yes </Question>
    <Question ID = 3.1> yes </Question>
    <Note> The person seems to be rather interested in the offered product </Note>
  </Survey>
</Phone campaign>
```



Custom scripts are still often required to perform JOIN operations or to enforce data consistency between documents



Document-oriented Databases

Documents and relationships

Relationships are represented by nesting documents: *"Data that is accessed together should be stored together"*

USERS						
ID	first_name	last_name	cell	city	latitude	longitude
1	Leslie	Yepp	8125552344	Pawnee	39.170344	-86.536632

HOBBIES		
ID	user_id	hobby
10	1	scrapbooking
11	1	eating waffles
12	1	working

JOB HISTORY			
ID	user_id	job_title	year_started
20	1	Deputy Directory	2004
21	1	City Councilor	2012
22	1	Director, National Parks Service, Midwest Branch	2014

```
{
  "_id": 1,
  "first_name": "Leslie",
  "last_name": "Yepp",
  "cell": "8125552344",
  "city": "Pawnee",
  "location": [ -86.536632,
                39.170344 ],
  "hobbies": [
    "scrapbooking",
    "eating waffles",
    "working"
  ],
  "jobHistory": [
    {
      "title": "Deputy Director",
      "yearStarted": 2004
    },
    {
      "title": "City Councilor",
      "yearStarted": 2012
    },
    {
      "title": "Director, National Parks
        Service, Midwest Branch",
      "yearStarted": 2014
    }
  ]
}
```



The core operations that a document-oriented database supports for documents are similar to other databases, and are named as CRUD:

- *Creation*: of a new document
- *Retrieval*: based on key, content, or metadata
- *Update*: the content or metadata of the document
- *Deletion*: of a document

Typical use cases: similar to key-value stores, but more complex data can be handled (e.g., product catalogues, user profiles).

→ **MongoDB** is an example of a document-oriented database

- It actually offers a very crude JOIN operation



The roots of column-oriented DBMSs can be traced back in the 1970s.


However, due to unfavorable market needs and non-easily accessible technology it was not until the 2000s that they took off.

A column-oriented DBMS stores table data physically by column instead of by row; in different disk locations, or even different machines.


Values belonging to the same column are packed together, as opposed to traditional DBMSs that store entire rows one after the other.



row-store



column-store





In a classical relational database each field is stored adjacent to the next in the same block on the hard drive:

```
512,Seabiscuit,Book,10.95,201712241200,goodreads.com
513,Bowler,Apparel,59.95,201712241200,google.com
514,Cuphead,Game,20.00,201712241201,gamerassaultweekly.com
```

In a columnar database, blocks on the disk for the above data might look like this:

```
512,513,514
Seabiscuit,Bowler,Cuphead
Book,Apparel,Game
10.95,59.95,20.00
201712241200,201712241200,201712241201
goodreads.com,google.com,gamerassaultweekly.com
```



Beneficial when the typical application is reading a subset of columns, or performing **aggregate functions** over them (AVG, MIN, MAX, ...).

Efficient storing capabilities due to an easier **compression of data**, which is performed by column (low information entropy).

Very scalable, and well suited for *massively parallel processing* (MPP), which involves having data partitioned and spread across a large cluster of machines.



The main drawbacks of column-based stores are related to tuple (re)construction overhead.

Newly inserted tuples have to be **broken down** to their component attributes, and each attribute is then written separately.

Also retrieving a tuple is considered problematic, since information about it may be stored in multiple locations, which brings an **overhead** for queries that access many attributes.



Column-oriented Stores

Difference wrt column-family stores


A column-oriented store *is not the same* as a column-family store! (and yes, computer scientists are hideous people)

A column-oriented database stores data tables by column rather than by row

- **Apache Druid**, **MonetDB** and **Vertica** are examples of column-oriented stores

A column-family database sees a row as a varying collection of columns, intuitively like a two-level key-value store (thus, it is simpler than a document store, but optimized for distributed scalability and huge datasets).


Cassandra, **Apache Hbase** and **Google BigTable** are examples of column-family stores (aka wide-column stores).





A graph database uses graph structures for semantic queries with *nodes* (possibly of different kinds), *edges*, and *properties* to represent and store data.

Graph databases can naturally represent certain kinds of semi-structured, **highly interconnected data**, such as those present in social networks, or in geospatial and biotech applications.





Of course, the previous graph could be modeled using other kinds of NoSQL approaches, as well as RDBMSs.

However, the resulting database would be difficult to query, update, and populate.

On the contrary, graph databases can easily answer to queries such as:

- what is the shortest path connecting node X with node Y?
- what are the friends of *Billy*?
- what are all, direct and indirect, friends of *Billy*?

In native implementations, nodes have **direct pointers** connecting them, so the navigation time is independent from the graph size.

- In other words, relationships are stored natively rather than simulated through joins as in happens with relational databases



Neo4j is a very performing, native, open source property graph database that guarantees ACID properties offering scalability up to billions of nodes.

- Online sandbox: <https://neo4j.com/sandbox/>


AllegroGraph is a closed source triplestore, designed to store RDF triples; it supports ACID properties.

ArangoDB is a multi-model, open source, database system that supports three data models (key/value, documents, graphs).

A graph database can be used to implement a so-called *Knowledge Graph* (KG), which is an example of **semantic network**.


A KG may enable:

- knowledge inference
- contextual reasoning
- complex relationship discovery



Knowledge graphs encode semantics **explicitly** through entities and relationships. Practical applications of knowledge graphs include:

- **Search engines:** understanding entities and their meanings
- **Recommendation systems:** connecting users, items, and preferences
- **Fraud detection:** finding suspicious connection patterns
- **Data integration:** linking heterogeneous data sources (ontology)
- **Question answering:** retrieving facts through semantic relations





Graph Databases

A relevant application: Knowledge graphs


Launched in 2012, the **Google Knowledge Graph** is a system that stores facts and information about entities assembled from materials shared across the web, as well as from open source and licensed databases.

It has amassed **over 500 billion facts** about five billion entities (as of May 2020).

It enables users to search for things, people, or places that Google knows about, instantly returning relevant information and helping to disambiguate between entities with similar names.

The screenshot displays the Google Knowledge Graph for Florence Price. Red arrows highlight several key features:

- Overview Tab:** The top navigation bar includes tabs for Overview, Videos, Songs, Albums, Listen, and People Also Search For.
- Available on:** A section showing where the user can listen to Florence Price's music, including Spotify, YouTube, Pandora, Apple Music, and Deezer.
- About:** A detailed biographical section containing information about her birth (April 9, 1887, Little Rock, AR), death (June 3, 1953, Chicago, IL), education (The University of Chicago, New England Conservatory of Music), children (Edith Cassandra Price, Thomas Jr., Florence Price Robinson), and albums (Symphonies: no. 1 in E minor / no. 4 in D minor, MORE).
- Songs:** A section listing specific works, such as "Dances in the Canebrakes" and "Fantasie Negre".




Modern AI systems based on deep learning represent data through **embeddings**.

An embedding is a numerical vector that **implicitly** captures the semantic meaning of an object, such as: text, images, audio, users or products.

E.g., the word “dog” might be represented by the embedding $[0.1, 2.5, \dots, -1.7]$

Objects with similar meanings are mapped to **nearby** points in a high-dimensional vector space.





The rise of AI applications based on embeddings has created the need for specialized databases to manage such high-dimensional vectors efficiently.

A vector database **stores, indexes, and retrieves embeddings** together with their associated metadata.


Its primary functionality is efficient **similarity search** over embeddings (e.g., nearest-neighbor queries).

Vector databases also support traditional filtering operations over metadata, which can be combined with embedding similarity search.

Popular vector databases: **Pinecone, Milvus, Weaviate, and Qdrant.**

Traditional databases search for exact matches expressed by queries (e.g., SQL).

A vector database, instead, given a query embedding, retrieves the nearest vectors according to a similarity metric, e.g., cosine similarity (approximate match).



Example

Query: “movies about space exploration”

Retrieved documents: “Interstellar”, “Apollo 13”, “The Martian”

- even if the exact words space exploration never appear within them



In recent years there has been a progressive **(re)approach between relational solutions and NoSQL**, with the birth of the so-called NewSQL

NewSQL DBMSs try to combine the performances of NoSQL with the data guarantees offered by relational DBs

For example, there are NoSQL systems that support transactions (Neo4j); others offer languages inspired by SQL, but generally less powerful than the latter (Cassandra's CQL)

Conversely, relational DBMSs allow, today, to store semi-structured data inside tables, thanks to the JSON and JSONB formats, while others allow to **store time series data** generated quickly and in large quantities (e.g., TimescaleDB), or to manage vectors (PostgreSQL's pgvector extension)



Stavros Harizopoulos, Daniel Abadi, Peter Boncz (2009), Column-Oriented Database Systems, VLDB 2009 Tutorial

Ian Robinson, Jim Webber & Emil Eifrem, Graph Databases Second Edition, O'Reilly Media, Inc.

Big Data Management and NoSQL Databases Lecture 7. Column-family stores, Doc. RNDr. Irena Holubova, Ph.D.

Andrew Pavlo, Matthew Aslett (2016), What's really new with NewSQL?, ACM Sigmod Record, 45(2), 45-55

Google Knowledge Graph: <https://blog.google/products-and-platforms/products/search/introducing-knowledge-graph-things-not/>

Vector Databases on Pinecone: <https://www.pinecone.io/learn/vector-database/>